

Package ‘BFI’

September 1, 2023

Type Package

Title Bayesian Federated Inference

Version 0.6.4

Date 2023-07-16

Author Hassan Pazira [aut, cre],
Marianne A. Jonker [aut],
Anthony C.C. Coolen [aut]

Maintainer Hassan Pazira <hassan.pazira@radboudumc.nl>

Description Bayesian Federated Inference method combines data from different (medical) centers without sharing them. In this version of the package, the user can fit models specifying Gaussian and Binomial (Logistic) families.

License GPL (>= 2)

URL <https://github.com/hassanpazira/BFI>

Encoding UTF-8

Suggests knitr,
rmarkdown,
testthat (>= 3.0.0)

VignetteBuilder knitr

Depends R (>= 2.10)

LazyData true

Imports devtools

Config/testthat/edition 3

RoxygenNote 7.2.3

R topics documented:

BFI-package	2
bfi	2
inv.prior.cov	5
MAP.estimation	7
Nurses	11
summary.bfi	11
trauma	14
Index	15

BFI-package

*Bayesian Federated Inference***Description**

Bayesian Federated Inference method combines inference results from different (medical) centers without sharing the data. In this version of the package, the user can fit models specifying Gaussian and Binomial (Logistic) families. The package will be updated with more models soon.

Details

Package: BFI
 Type: Package
 Version: 0.6.4
 Date: 2023-07-16
 License: GPL (>=2)

MAP.estimation and bfi are the main functions. All other functions are utility functions.

Some examples are provided in the vignettes accompanying this package in order to show how the package can be applied to a real data.

Author(s)

Hassan Pazira, Marianne A. Jonker, Anthony C.C. Coolen
 Maintainer: Hassan Pazira <hassan.pazira@radboudumc.nl>

References

Jonker M.A., Pazira H. and Coolen A.C.C. (2023). *Bayesian Federated Inference for Statistical Models, Statistics in Medicine*, Vol. 0(0), 0-0. <<https://doi.org/10.48550/arXiv.2302.07677>>

bfi

*Bayesian Federated Inference***Description**

bfi function can be used (in central server) to combine inference results from separate data sets (without combining the data) to approximate what would have been inferred had the data sets been merged. For now the function can handle linear and logistic regression models, but more models will be available in the near future. bfi command

Usage

```
bfi(theta_hats = NULL, A_hats, Lambda, const_var = NULL,
    stratified = FALSE, nuisance = 1L)
```

Arguments

theta_hats	a list of L vectors of the maximum a posteriori (MAP) estimates of the model parameters in the L centers. These vectors must have equal dimensions. See 'Details'.
A_hats	a list of L curvature matrices for L centers. These matrices must have equal dimensions. See 'Details'.
Lambda	a list of $L + 1$ matrices. Each matrix used as the prior of the inverse variance-covariance matrix of Gaussian distribution. The first L matrices are for L local centers and the last one is the chosen variance-covariance matrix for the Gaussian prior of the (fictive) combined data set. These matrices must have equal dimensions. If all the $L + 1$ matrices are the same, this argument can be one matrix or a list of one matrix. See 'Details'.
const_var	a vector of L elements. Each element represents a specific feature of the corresponding center. There must be only one specific value or attribute for each center. This vector could be a numeric, characteristic or factor vector.
stratified	logical flag for performing the stratified analysis. If stratified=TRUE, the parameters selected in the nuisance argument can be different across centers. Default is stratified=FALSE. See 'Details'.
nuisance	a vector of one or two elements. It is usable if stratified=TRUE. For the binomial family the length of the vector should be one which refers to 'intercept', and for gaussian this vector can have maximum two elements which refer to 'intercept' and/or 'sigma2'. See 'Details'.

Details

bfi function implements the BFI approach described in the paper Jonker et. al. (2023) given in the references. The results gathered from different (L) centers are combined, and the BFI estimates of the model parameters and curvature matrix evaluated at that point are returned.

The result from each center must be obtained using the MAP.estimation function separately, and then all of these results (coming from different centers) should be compiled into a list to be used as an input of bfi(). The models in the different centers should be defined in exactly the same way; among others, exactly the same covariates should be included in the models. The parameter vectors should be defined exactly the same, so that the estimates theta_hat's and A_hat's (the output of the MAP.estimation function) are defined in the same way.

Note that, the order of the estimates in the lists theta_hats, A_hats and Lambda must be the same, so that in every list the elements at the ℓ position are all from the same center. For example, the first three entries of bfi() could be as follows:

```
bfi(theta_hats=theta_list, A_hats=A_list, Lambda= Lambda_list)
```

where

```
theta_list=list(obj1$theta_hat, ..., obj_L$theta_hat, ..., obj_L$theta_hat),
```

```
A_list=list(obj1$A_hat,..., obj_L$A_hat,..., obj_L$A_hat),
```

```
Lambda_list=list(obj1$Lambda,..., obj_L$Lambda,..., obj_L$Lambda),
```

and

```
obj_L=MAP.estimation(y=y_L, X=X_L, Lambda=Lambda_L) # must be done at the center  $\ell$ 
```

where y_L , X_L and $Lambda_L$ are the data and the prior of the inverse variance-covariance matrix for the ℓ center.

Details related to stratified=TRUE should be added!!!!...

Details related to const_var should be added!!!!...

Value

bfi returns a list containing the following components:

theta_hat	the p - or $(p + 1)$ -dimensional vector of estimates obtained by combining the inference results from the L centers with the 'BFI' methodology when stratified=FALSE, where p is the number of regression parameters including intercept. If stratified=TRUE, dimension of this vector is $L + p - 1$ (for binomial family, or for gaussian family when nuisance has one element) or $2 * L + p - 1$ (for gaussian family when nuisance has two elements). See 'Details';
A_hat	the curvature matrix obtained by the 'BFI' method when stratified=FALSE, such that it's a $(p \times p)$ - or $(p + 1) \times (p + 1)$ -dimensional matrix depending on the family used, where p is the number of regression parameters including intercept. If stratified=TRUE, it's a list of L matrices corresponding to each center. This is not the same as A_hats in 'Arguments'. See 'Details';
sd	the p - or $(p + 1)$ -dimensional vector of standard deviation of the estimates in theta_hat if stratified=FALSE, i.e., the vector equals $\sqrt{\text{diag}(\text{solve}(\text{A_hat}))}$. If stratified=TRUE, it's a list of L vectors; one vector for every center. These vectors are standard deviation of parameter estimates obtained from the matrices in A_hat, i.e., the vector equals $\sqrt{\text{diag}(\text{solve}(\text{A_hat}[[j]]))}$ where j refers to a center.

Author(s)

Hassan Pazira

Maintainer: Hassan Pazira <hassan.pazira@radboudumc.nl>

References

Jonker M.A., Pazira H. and Coolen A.C.C. (2023). *Bayesian Federated Inference for Statistical Models*. Statistics in Medicine, Vol. 0(0), 0-0. <<https://doi.org/10.48550/arXiv.2302.07677>>

See Also

[MAP.estimation](#)

Examples

```
#-----
# y ~ Binomial
#-----
set.seed(112358)
p <- 4 # number of regression coefficients (including intercept), is the same for all L=2 locations

##-----
## Local center 1
##-----
n1 <- 30
X1 <- data.frame(x1=rnorm(n1), # standard normal variable
                 x2=sample(0:1, n1, replace=TRUE), # dichotomous variable
                 x3=sample(1:3, n1, replace=TRUE)) # categorical variable
#true_beta <- rep(2, p) # with an intercept b0=b1=b2=b3=2
eta1 <- 2 + apply(2 * X1, 1, sum) # linear predictor: X%*\beta=\eta
mu1 <- binomial()$linkinv(eta1) # inverse of the link function: g^{-1}(\eta)=\mu
y1 <- rbinom(n1, 1, mu1)
```

```

# we assume the same inverse covariance matrix for all locations
Lambda      <- inv.prior.cov(X1, 0.01, family=binomial)
fit1        <- MAP.estimate(y1, X1, family=binomial, Lambda)
theta_hat1  <- fit1$theta_hat # intercept and coefficient estimates
A_hat1      <- fit1$A_hat     # curvature matrix

##-----
## Local center 2
##-----
n2          <- 50
X2          <- data.frame(x1=rnorm(n2),                # standard normal variable
                          x2=sample(0:1, n2, replace=TRUE), # dichotomous variable
                          x3=sample(1:3, n2, replace=TRUE)) # categorical variable
eta2        <- 2 + apply(2 * X2, 1, sum) # linear predictor: X%*\beta=\eta
mu2         <- binomial()$linkinv(eta2) # inverse of the link function: g^{-1}(\eta)=\mu
y2          <- rbinom(n2, 1, mu2)
fit2        <- MAP.estimate(y2, X2, family=binomial, Lambda)
theta_hat2  <- fit2$theta_hat
A_hat2      <- fit2$A_hat

###-----
### Bayesian Federated Inference
###-----
A_hats      <- list(A_hat1, A_hat2)
theta_hats  <- list(theta_hat1, theta_hat2)
bfi(theta_hats, A_hats, Lambda)

```

inv.prior.cov

Creates an inverse covariance matrix for a Gaussian prior

Description

inv.prior.cov builds a diagonal inverse covariance matrix for the Gaussian prior distribution based on the design matrix of covariates, that takes into account the number of regression parameters in case of categorical covariates. In case of a linear model, it also includes the variance of the measurement errors.

Usage

```
inv.prior.cov(X, lambda = 1, family = gaussian, intercept = TRUE,
              independ = TRUE, set_seed = NULL)
```

Arguments

X	design matrix of dimension $n \times p$, where p is the number of covariates/predictors without intercept.
lambda	the tuning parameter for ridge (L2) penalization that controls the strength of the ridge penalty. It could be either a single positive number, a vector of two elements, or a vector with length equal to p , $p+1$ or $p+2$ (where p is the number of covariates, i.e., number of columns of X argument) depends on the arguments X, family, and intercept. If lambda chosen by the user is a single positive number, the function inv.prior.cov() consider lambda as a vector whose all

	elements are equal to that scalar value. Default is <code>lambda=1</code> . If the user choose a vector of only two elements as an entry value, the <code>inv.prior.cov()</code> function set <code>lambda</code> as a vector whose all elements, except the last one, are equal to the first element of the entry, and the last element is set to the second element of the entry. Used when <code>independ=TRUE</code> .
<code>family</code>	a description of the error distribution and link function used to specify the model. This can be a character string naming a family function or the result of a call to a family function (see family for details). By default the gaussian family (with identity link function) is used.
<code>intercept</code>	logical flag for fitting an intercept. The intercept is fitted if <code>intercept=TRUE</code> (the default) or set to zero if <code>intercept=FALSE</code> . If <code>intercept=FALSE</code> , the first row and column of this matrix are set to zero
<code>independ</code>	logical flag for creating the prior covariance matrix with (in)dependent variables/parameters. If <code>TRUE</code> (the default), the result is a diagonal matrix which means the parameters are independent. If <code>FALSE</code> , the off-diagonal elements of the resulted matrix are generated from the standard normal distribution <code>rnorm()</code> , and the diagonal elements are the vector <code>lambda</code> .
<code>set_seed</code>	if <code>independ=FALSE</code> , this argument (which should be an integer) ensures consistent results. If <code>set_seed=NULL</code> (the default), the results are not consistent.

Details

`inv.prior.cov` creates a matrix whose dimension depends on the arguments `X`, `family`, and `intercept`. If we assume the parameters are independent (`independ=TRUE`), the function `inv.prior.cov` returns a diagonal matrix with the vector `lambda` as its diagonal. If `independ=FALSE`, `inv.prior.cov` returns a matrix which is equal to $z \cdot z^t$ where z is a matrix such that all elements are generated by standard normal distribution.

The output of `inv.prior.cov()` is used in the main functions `MAP.estimation()` and `bfi()`.

Value

`inv.prior.cov` returns a matrix. The dimension of the matrix depends on the arguments `X`, `family`, and `intercept`.

Author(s)

Hassan Pazira
Maintainer: Hassan Pazira <hassan.pazira@radboudumc.nl>

References

Jonker M.A., Pazira H. and Coolen A.C.C. (2023). *Bayesian Federated Inference for Statistical Models*, Statistics in Medicine, Vol. 0(0), 0-0. <<https://doi.org/10.48550/arXiv.2302.07677>>

See Also

[MAP.estimation](#)

Examples

```
#-----
# y ~ Binomial
#-----
X      <- data.frame(matrix(rnorm(50 * 4), 50, 4))
lambda <- 0.05
# We assume the same (inverse) covariance matrix for all
# locations and consider independency of prior parameters
(Lambda <- inv.prior.cov(X, lambda, family=binomial))
# No intercept
(Lambda <- inv.prior.cov(X, lambda, family="binomial", intercept = FALSE))

#-----
# y ~ Gaussian
#-----
X      <- data.frame(matrix(rnorm(50 * 3), 50, 3))
lambda <- 0.01
sigma2e <- 0.5
# If we consider dependency for all prior parameters:
(Lambda <- inv.prior.cov(X, family="gaussian", independ = FALSE, set_seed=1123))
# No intercept
(Lambda <- inv.prior.cov(X, family="gaussian", intercept = FALSE, independ = FALSE, set_seed=1123))
```

MAP.estimation

Maximum A Posteriori estimation

Description

MAP.estimation function is used (in local centers) to estimate Maximum A Posterior (MAP) of the parameters for the GLM and Survival models.

Usage

```
MAP.estimation(y, X, family = gaussian, Lambda, intercept = TRUE,
               initial = NULL, control = list())
```

Arguments

- | | |
|--------|---|
| y | response vector. If the binomial family is used, this argument can be a vector with entries 0 (failure) or 1 (success). Alternatively, the response can be a matrix where the first column is the number of “successes” and the second column is the number of “failures”. |
| X | design matrix of dimension $n \times p$, where p is the number of covariates (predictors) plus intercept. |
| family | a description of the error distribution and link function used to specify the model. This can be a character string naming a family function or the result of a call to a family function (see family for details). By default the gaussian family (with identity link function) is used. |
| Lambda | the matrix used as the prior of the inverse variance-covariance matrix of Gaussian distribution. It can be created by the function <code>inv.prior.cov()</code> . |

intercept	logical flag for fitting an intercept. The intercept is fitted if intercept=TRUE (the default) or set to zero if intercept=FALSE. Although the intercept is included in the model by default, it can be penalized.
initial	a vector specifying initial values for the parameters (intercept, coefficients and/or error variance) to be optimized over. For the gaussian family, it should be a $p + 1$ -dimensional vector with a non-negative value for the last element, and for binomial the length of the vector should be p , where p is the number of covariates plus intercept. Since the 'L-BFGS-B' method is used in the algorithm, these values should always be finite. Default is a vector of zeros.
control	a list of control parameters. See 'Details'.

Details

MAP.estimation function finds the Maximum A Posteriori (MAP) estimates of model parameters, i.e., the values associated with the peak of the log-posterior distribution (the posterior mode). In other words, MAP.estimation() optimizes the log-posterior density with respect to the parameter vector to obtain its MAP estimation. In addition to the model parameters, i.e., coefficients (β 's) and variance error (σ_e^2), the curvature matrix (Hessian of the log-posterior) is estimated around the mode.

The current version of the **BFI** package supports two families of GLM: gaussian and binomial. For the gaussian family, MAP.estimation function returns the $(p + 1)$ -dimensional vector of the MAP estimates of the coefficients $(\beta_0, \beta_1, \dots, \beta_{p-1})$ and variance error parameter, respectively. When the binomial family is used, it returns the MAP estimates of the p coefficients while the variance error or dispersion is set to one.

Note that: although an intercept is included in the model by default (except for family=cox) using intercept=TRUE, it can be penalized. It can also be set to zero using intercept=FALSE.

Since all elements of the lists in the first three arguments of the bfi function must have equal dimensions, the function MAP.estimation creates the output for all centers with similar dimensions even if in some centers the intercept is set to zero intercept=FALSE (of course for the same family's) by setting the corresponding elements to zero.

The MAP.estimation function returns an object of class 'bfi'. Therefore, summary() can be used for the object returned by MAP.estimation().

To solve unconstrained and bound-constrained optimization problems, the MAP.estimation function utilizes an optimization algorithm called Limited-memory Broyden-Fletcher-Goldfarb-Shanno with Bound Constraints (L-BFGS-B), Byrd et. al. (1995). The L-BFGS-B algorithm is a limited-memory "quasi-Newton" method that iteratively updates the parameter estimates by approximating the inverse Hessian matrix using gradient information from the history of previous iterations. This approach allows the algorithm to approximate the curvature of the posterior distribution and efficiently search for the optimal solution, which makes it computationally efficient for problems with a large number of variables.

By default, the algorithm uses a relative change in the objective function as the convergence criterion. When the change in the objective function between iterations falls below a certain threshold ('factr') the algorithm is considered to have converged.

In Jonker et. al. (2023), the Gaussian and non-informative priors were used on the parameters. In this case, the posterior distribution is highly expected to be unimodal for a Gaussian GLM, which means it has only one minimum and no local optima. This property ensures that the optimization algorithm will converge to the global minimum, guaranteeing convergence. After fitting the model, the convergence status can be examined by checking the convergence attribute of the model object using the \$convergence attribute. See 'Value' for the possible options of this attribute.

If a logistic regression model (`family=binomial`) is used, it's generally more likely that the posterior distribution will be unimodal with these two priors, especially if the data provides strong information about the parameters or sample size is large. Since the starting point has a great impact on the converged point especially for densities which are not unimodal, multiple starting points can be used in the `initial` argument to check for unimodality. Indeed, the negative log-posterior density is minimized from an chosen starting point `initial`.

However, there could be cases with small data sets (non-informative data), high-dimensionality, or inappropriate optimization settings, which might lead to multimodality or other complex behavior in the posterior distribution, resulting in users experiencing convergence issues. In such cases, it may be necessary to investigate and adjust optimization parameters to facilitate convergence. It can be done using the `control` argument.

The argument `control` is a list that can supply any of the following components:

`maxit`: is the maximum number of iterations. Default is `1e2`;

`factr`: controls the convergence of the 'L-BFGS-B' method. Convergence occurs when the reduction in the objective is within this factor of the machine tolerance. Default is `1e7`, that is a tolerance of about `1e-8`;

`pgtol`: helps control the convergence of the 'L-BFGS-B' method. It is a tolerance on the projected gradient in the current search direction. Default is zero, when the check is suppressed;

`trace`: is a non-negative integer. If positive, tracing information on the progress of the optimization is produced. Higher values may produce more tracing information: for the method 'L-BFGS-B' there are six levels of tracing. To understand exactly what these do see the source code of `optim` function in the [stats](#) package;

`REPORT`: is the frequency of reports for the 'L-BFGS-B' method if '`control$trace`' is positive. Default is every 10 iterations;

`lmm`: is an integer giving the number of BFGS updates retained in the 'L-BFGS-B' method. Default is 5.

Value

`MAP.estimation` returns a list containing the following components:

<code>theta_hat</code>	the p - or $(p + 1)$ -dimensional vector corresponding to the maximum a posteriori (MAP) estimation of the parameters, where p is the number of covariates plus intercept. If <code>intercept=FALSE</code> , the first element of the vector is set to zero;
<code>A_hat</code>	the curvature matrix around the point <code>theta_hat</code> . It's a $(p \times p)$ - or $(p + 1) \times (p + 1)$ -dimensional matrix depending on the family used. If <code>intercept=FALSE</code> , the first row and column of the matrix are set to zero;
<code>sd</code>	the p - or $(p + 1)$ -dimensional vector of standard deviation of estimates in <code>theta_hat</code> , i.e., <code>sqrt(diag(solve(A_hat)))</code> . If <code>intercept=FALSE</code> , the first element of this vector is set to zero;
<code>Lambda</code>	the matrix used as the prior of the inverse variance-covariance matrix. If <code>intercept=FALSE</code> , the first row and column of this matrix are set to zero;
<code>formula</code>	the formula of the model;
<code>n</code>	sample size;
<code>np</code>	the number of coefficients/regression parameters. In other words, number of predictors plus the intercept if <code>intercept=TRUE</code> , or without intercept if <code>intercept=FALSE</code> ;
<code>value</code>	the value of the 'negative' log-likelihood function corresponding to <code>theta_hat</code> ;
<code>family</code>	a description of the error distribution used in the model;

intercept	logical flag used to fit an intercept if TRUE, or set to zero if FALSE;
convergence	an integer value used to encode the warnings and the errors related to the algorithm used to fit the model. The values returned are: 0 algorithm has converged; 1 maximum number of iterations ('maxit') has been reached; 2 Warning from the 'L-BFGS-B' method. See the message after this value;
control	the list of control parameters used to compute the MAP estimates.

Author(s)

Hassan Pazira

Maintainer: Hassan Pazira <hassan.pazira@radboudumc.nl>

References

Jonker M.A., Pazira H. and Coolen A.C.C. (2023). *Bayesian Federated Inference for Statistical Models*. Statistics in Medicine, Vol. 0(0), 0-0. <<https://doi.org/10.48550/arXiv.2302.07677>>

Byrd R.H., Lu P., Nocedal J. and Zhu C. (1995). *A limited memory algorithm for bound constrained optimization*. SIAM Journal on Scientific Computing, 16, 1190-1208. <<https://doi.org/10.1137/0916069>>

See Also

[bfi](#) and [summary.bfi](#)

Examples

```
#-----
# y ~ Gaussian
#-----

set.seed(11235813)
n <- 30
p <- 3 # number of coefficients (without intercept)
X <- data.frame(matrix(rnorm(n * p), n, p))
eta <- 1 + 2 * X[,1] # with an intercept
mu <- gaussian()$linkinv(eta)
sigma2 <- 1.5
# the true theta is c(1, 2, 0, 0, sigma2)
y <- rnorm(n, mu, sd = sqrt(sigma2))
lambda <- 0.01
# inverse of covariance matrix:
Lambda <- inv.prior.cov(X, lambda = c(lambda, sigma2), family = gaussian)

# MAP estimates of the parameters of interest (including 'intercept') and curvature matrix
(fit <- MAP.estimation(y, X, family = gaussian, Lambda))
class(fit)

# MAP estimates without 'intercept'
Lambda <- inv.prior.cov(X, lambda = c(lambda, sigma2), family = gaussian, intercept = FALSE)
(fit1 <- MAP.estimation(y, X, family = gaussian, Lambda, intercept = FALSE))
```

Nurses

*Nurses' stress in different hospitals***Description**

This data set contains three-level simulated data from a hypothetical study on stress in hospitals. The data are from nurses working in wards nested within hospitals. It is a cluster-randomized experiment. In each of 25 hospitals, four wards are selected and randomly assigned to an experimental and a control condition. In the experimental condition, a training program is offered to all nurses to cope with job-related stress. After the program is completed, a sample of about 10 nurses from each ward is given a test that measures job-related stress. Additional variables are: nurse age (years), nurse experience (years), nurse gender (0 = male, 1 = female), type of ward (0 = general care, 1 = special care), and hospital size (0 = small, 1 = medium, 2 = large).

Usage

```
data(Nurses)
```

Source

<https://multilevel-analysis.sites.uu.nl/datasets/>

References

Hox, J., Moerbeek, M., and van de Schoot, R. (2010). *Multilevel Analysis: Techniques and Applications*, Second Edition (2nd ed.). Routledge. <<https://doi.org/10.4324/9780203852279>>

summary.bfi

*Summarizing BFI Fits***Description**

Summary method for an object with class 'bfi' created by the `MAP.estimation` function.

Usage

```
## S3 method for class 'bfi'
summary(object, curmat = FALSE,
        digits = max(3, getOption("digits") - 3), ...)
```

Arguments

<code>object</code>	fitted <code>bfi</code> object.
<code>curmat</code>	logical; if TRUE, the curvature matrix around the estimated parameters is returned and printed. Default is FALSE.
<code>digits</code>	significant digits in printout.
<code>...</code>	additional arguments affecting the summary produced.

Details

`summary.bfi()` gives information about the MAP estimates of parameters of the model. It can be used for the `bfi` objects built by the `MAP.estimate` function.

The output of the summary method shows the details of the model, i.e. formula, family and link function used to specify the generalized linear model, followed by information about the estimates, standard deviations and confidence intervals. Information about the log-likelihood and convergence status are also provided.

For a model fitted without an intercept (i.e., in `MAP.estimate()` the argument is `intercept=FALSE`), `summary.bfi()` exhibits the outputs without the intercept unlike the outputs of `MAP.estimate()` which are replaced by zeros.

By default, `summary.bfi` function does not return the curvature matrix, but the user can use `curmat=TRUE` to print it.

Value

`summary.bfi` returns an object of class `summary.bfi`, a list with the following components:

<code>theta_hat</code>	the component from object. The last element of this vector is the estimate of the dispersion parameter (<code>sigma2</code>). See the MAP.estimate function.
<code>A_hat</code>	the component from object. See the MAP.estimate function.
<code>sd</code>	the component from object. The last element of this vector is the square root of the estimated dispersion. See the MAP.estimate function.
<code>Lambda</code>	the component from object. See the MAP.estimate function.
<code>formula</code>	the component from object. See the MAP.estimate function.
<code>n</code>	the component from object. See the MAP.estimate function.
<code>np</code>	the component from object. See the MAP.estimate function.
<code>value</code>	the component from object. See the MAP.estimate function.
<code>family</code>	the component from object. See the MAP.estimate function.
<code>intercept</code>	the component from object. See the MAP.estimate function.
<code>convergence</code>	the component from object. See the MAP.estimate function.
<code>control</code>	the component from object. See the MAP.estimate function.
<code>Estimate</code>	the coefficients without the estimate <code>sigma2</code> . Moreover, if the intercept was not fitted (i.e., <code>object\$intercept</code> is <code>FALSE</code>), <code>(Intercept)</code> is not shown here, while it can be seen in <code>object\$theta_hat</code> which is zero. See 'Examples'.
<code>se</code>	the standard error. Here <code>se</code> is equal to the standard deviation (<code>sd</code>).
<code>CurMat</code>	the curvature matrix. If the intercept was not fitted (i.e., <code>object\$intercept</code> is <code>FALSE</code>), the row and column related to the intercept will not be shown here unlike <code>object\$A_hat</code> . See 'Examples'.
<code>logLik</code>	the value of the loglikelihood function corresponding to estimates (<code>theta_hat</code>). This is the minus of the value component.
<code>link</code>	the link function. By default the gaussian family with identity link function and the binomial family with logit link function are used.
<code>dispersion</code>	the estimated variance of the random error, i.e., <code>sigma2</code> . The dispersion is taken as 1 for the binomial family.
<code>CI</code>	a 95% confidence interval.

Author(s)

Hassan Pazira

Maintainer: Hassan Pazira <hassan.pazira@radboudumc.nl>

See Also[MAP.estimation](#) and [bfi](#)**Examples**

```

#-----
# y ~ Gaussian (with intercept)
#-----

set.seed(1123581)
n      <- 30
p      <- 4
X      <- data.frame(matrix(rnorm(n * p), n, p))
b      <- 1:2
eta    <- b[1] + X[, 1] * b[2]
mu     <- gaussian()$linkinv(eta)
sigma2e <- 0.5
y      <- rnorm(n, mu, sd=sqrt(sigma2e))
lambda <- 0.1
Lambda <- inv.prior.cov(X, lambda=c(lambda,sigma2e), family=gaussian)
fit    <- MAP.estimation(y, X, family=gaussian, Lambda)
class(fit)

summary(fit)
sumfit <- summary(fit, curmat = TRUE)
sumfit$logLik
sumfit$dispersion
sumfit$CI
class(sumfit)

#-----
# y ~ Binomial (without intercept)
#-----

n0      <- 40
p0      <- 5
X0      <- data.frame(matrix(rnorm(n0 * p0), n0, p0))
eta0    <- 1 + X0[, 1] * 2
mu0     <- binomial()$linkinv(eta0)
y0      <- rbinom(n0, 1, mu0)
Lambda0 <- inv.prior.cov(X0, lambda=0.5, family=binomial, intercept=FALSE)
fit0    <- MAP.estimation(y0, X0, family=binomial, Lambda0, intercept=FALSE)

sumfit0 <- summary(fit0, curmat = TRUE)
sumfit0$A_hat
sumfit0$theta_hat[-length(sumfit0$theta_hat)]
sumfit0$Estimate

```

trauma*Trauma patients from different hospitals*

Description

This data set consists of data of 371 trauma patients from three hospitals. The binary variable mortality is used as an outcome, and variables age, sex, the Injury Severity Score (ISS, ranging from 1 (low) to 75 (high)) and the Glasgow Coma Scale (GCS, which expresses the level of consciousness, ranging from 3 (low) to 15 (high)) are used as covariates. The data originate from multiple hospitals which can be categorized in three groups as: peripheral hospital without a neuro-surgical unit (Status = 1), peripheral hospital with a neuro-surgical unit (Status = 2), and academic medical center (Status = 3).

Usage

```
data(trauma)
```

References

- Jonker M.A., Pazira H. and Coolen A.C.C. (2023). *Bayesian Federated Inference for Statistical Models*, *Statistics in Medicine*, Vol. 0(0), 0-0. <<https://doi.org/10.48550/arXiv.2302.07677>>
- Draaisma J.M.Th, de Haan A.F.J., Goris R.J.A. (1989). *Preventable Trauma Deaths in the Netherlands - A prospective Multicentre Study*, *The journal of Trauma*, Vol. 29(11), 1552-1557.

Index

* **Bayesian**

- bfi, [2](#)
- BFI-package, [2](#)
- inv.prior.cov, [5](#)
- MAP.estimation, [7](#)
- summary.bfi, [11](#)

* **Federated**

- bfi, [2](#)
- BFI-package, [2](#)
- MAP.estimation, [7](#)
- summary.bfi, [11](#)

* **datasets**

- Nurses, [11](#)
- trauma, [14](#)

* **package**

- BFI-package, [2](#)

bfi, [2](#), [10](#), [13](#)

BFI-package, [2](#)

family, [6](#), [7](#)

inv.prior.cov, [5](#)

MAP.estimation, [4](#), [6](#), [7](#), [12](#), [13](#)

Nurses, [11](#)

stats, [9](#)

summary(summary.bfi), [11](#)

summary.bfi, [10](#), [11](#)

trauma, [14](#)